

Nithin S
221IT085

IT464 Lab Assignment 1

Q1. Find the line of best fit (linear regression) for the following data set and plot it. <https://www.kaggle.com/datasets/tanuprabhu/linear-regression-dataset> Find the predictions to $x=[20,40,60,...,280,300]$ and Compute the least squared error (LSE). Print/tabulate "x, predictions and LSE".

Code

```
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
import matplotlib.pyplot as plt

df = pd.read_csv('LR.csv')
X = df['X'].values.reshape(-1, 1)
y = df['Y'].values

model = LinearRegression()
model.fit(X, y)

x_pred = np.arange(20, 301, 20).reshape(-1, 1)
y_pred = model.predict(x_pred)

y_pred_original = model.predict(X)

# Calculate the total Least Squared Error (LSE) for the original data
lse_total = np.sum((y - y_pred_original) ** 2)

# Use the mean of the actual y values as the "actual" values for the predicted dataset
y_mean = np.mean(y)
y_actual_pred = np.full_like(y_pred, y_mean)

lse_individual_pred = (y_actual_pred - y_pred) ** 2

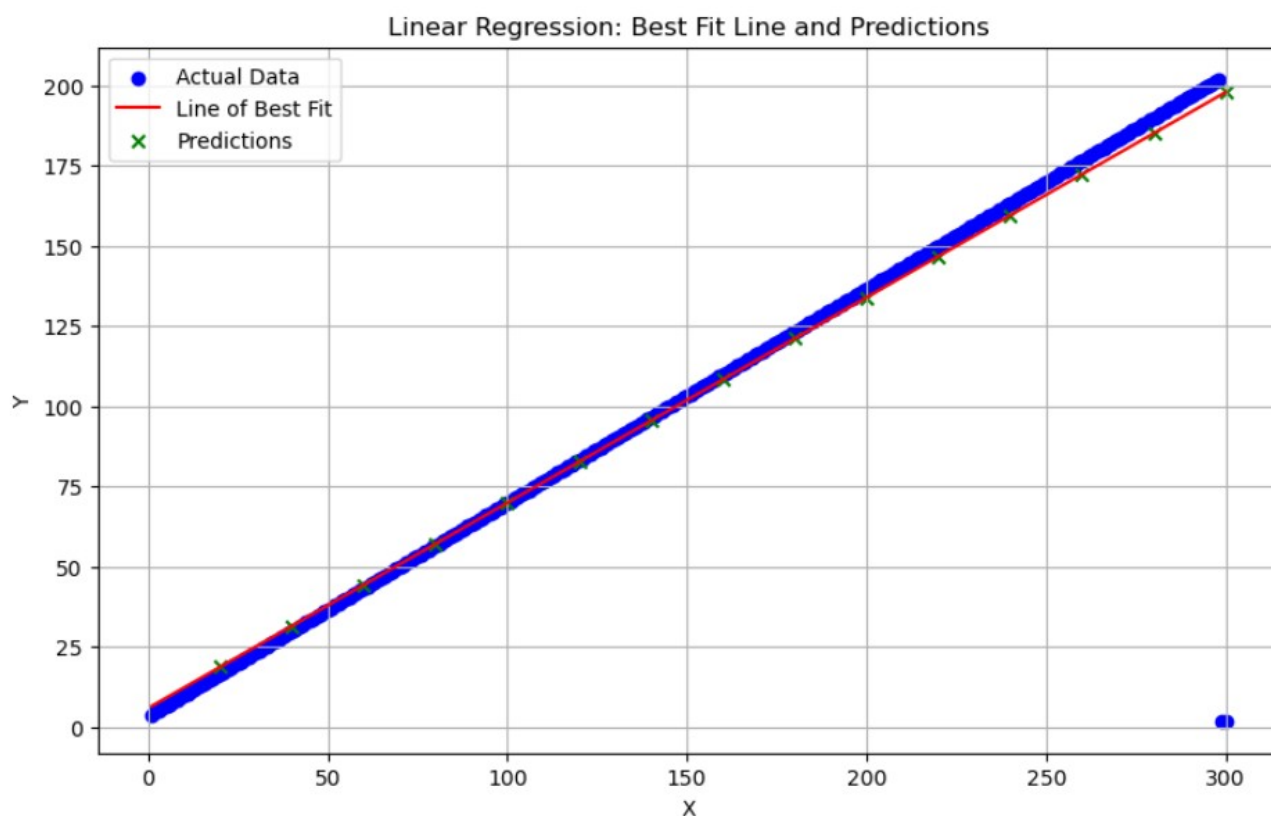
r2 = r2_score(y, y_pred_original)
accuracy_percentage = r2 * 100

results_df = pd.DataFrame({
    'X': x_pred.flatten(),
    'Predicted Y': y_pred,
    'Actual Y': y_actual_pred,
    'Individual LSE': lse_individual_pred
})

plt.figure(figsize=(10, 6))
plt.scatter(X, y, color='blue', label='Actual Data')
plt.plot(X, model.predict(X), color='red', label='Line of Best Fit')
plt.scatter(x_pred, y_pred, color='green', marker='x', label='Predictions')
plt.xlabel('X')
plt.ylabel('Y')
plt.title('Linear Regression: Best Fit Line and Predictions')
plt.legend()
plt.grid(True)
plt.show()

# Print the regression equation, LSE, and R-squared score
print(f"\nRegression Equation: y = {model.coef_[0]:.4f}x + {model.intercept_:.4f}")
print(f"\nTotal Least Squared Error (LSE): {lse_total:.4f}")
print(f"\nR-squared Score (R²): {r2:.4f}")
print(f"\nAccuracy (R² in percentage): {accuracy_percentage:.2f}%")

# Display results
print("\nPredictions and LSE for Predicted Data:")
print(results_df.to_string(index=False))
```



Regression Equation: $y = 0.6400x + 5.8888$
Total Least Squared Error (LSE): 78668.9421
R-squared Score (R^2): 0.9214
Accuracy (R^2 in percentage): 92.14%

Predictions and LSE for Predicted Data:

X	Predicted_Y	Individual_LSE
20	18.689697	6976.568967
40	31.490595	5002.020007
60	44.291493	3355.197016
80	57.092391	2036.099993
100	69.893289	1044.728939
120	82.694186	381.083853
140	95.495084	45.164735
160	108.295982	36.971586
180	121.096880	356.504405
200	133.897778	1003.763193
220	146.698675	1978.747949
240	159.499573	3281.458673
260	172.300471	4911.895366
280	185.101369	6870.058027
300	197.902267	9155.946656

Q2. Find the line of best fit (multiple linear regression - MLR) for the California Housing Dataset and plot it.
<https://www.geeksforgeeks.org/dataset-for-linear-regression/>
Note: Exclude “longitude, latitude and ocean proximity” parameters/variables.

Code

```
import pandas as pd
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import numpy as np

housing_data = pd.read_csv('housing.csv')
housing_data = housing_data.drop(['longitude', 'latitude', 'ocean_proximity'], axis=1)

X = housing_data.drop('median_house_value', axis=1)
y = housing_data['median_house_value']

imputer = SimpleImputer(strategy='median')
X_imputed = imputer.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_imputed, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

mlr_model = LinearRegression()
mlr_model.fit(X_train_scaled, y_train)

print("Coefficients:", mlr_model.coef_)
print("Intercept:", mlr_model.intercept_)
y_pred = mlr_model.predict(X_test_scaled)
plt.scatter(y_test, y_pred)
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Actual vs Predicted Prices')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red') # Line of best fit
plt.show()


import pandas as pd

housing_test_data = pd.read_csv("housing2.csv", skiprows=2)
housing_test_data = housing_test_data.drop(housing_test_data.columns[0], axis=1)
housing_test_data = housing_test_data.drop(index=housing_test_data.index[0]).reset_index(drop=True)

housing_test_data.columns
housing_test_data

expected_columns = ['housing_median_age', 'total_rooms', 'total_bedrooms', 'population', 'households', 'median_income']
housing_test_data = housing_test_data[expected_columns]
housing_test_data.dropna()
len(housing_test_data)
housing_test_data_imputed = imputer.transform(housing_test_data)
housing_test_data_scaled = scaler.transform(housing_test_data_imputed)
len(housing_test_data_scaled)
housing_test_predictions = mlr_model.predict(housing_test_data_scaled)
housing_test_data['Predicted_Price'] = housing_test_predictions
housing_test_data
```

Coefficients: [23906.19520346 -43497.74884467 42713.1788278 -40536.87160442
48392.8781107 91325.09220035]
Intercept: 207194.6937378882



Mean Absolute Error (MAE): 56713.67
Mean Squared Error (MSE): 5968852333.91
Root Mean Squared Error (RMSE): 77258.35
R-squared Score (R^2): 0.5445

Test Data with Predictions:

[11]:	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	Predicted_Price
0	20.0	4113.0	302.0	999.0	302.0	6.7905	266969.176243
1	10.0	7837.0	645.0	2307.0	265.0	2.0460	-70398.852854
2	63.0	8677.0	619.0	1659.0	227.0	8.5273	339800.487261
3	46.0	3754.0	984.0	607.0	334.0	6.6504	404318.088162
4	10.0	4301.0	578.0	653.0	391.0	5.5973	238791.306407
5	23.0	4052.0	458.0	1458.0	285.0	6.7502	269326.887482
6	9.0	9692.0	791.0	1472.0	369.0	5.3007	104554.690633
7	42.0	4602.0	916.0	657.0	353.0	7.8508	431036.493171
8	5.0	7083.0	101.0	1289.0	383.0	8.4030	235899.608621
9	33.0	5450.0	404.0	980.0	481.0	3.7130	151106.685909

Q3. Perform MLR and Logistic regression on the following data to predict heart disease. <https://www.kaggle.com/datasets/dileep070/heart-disease-prediction-using-logistic-regression> Predict heart disease for the “heart2” test data.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.metrics import (ConfusionMatrixDisplay, classification_report, accuracy_score,
                             roc_curve, auc, mean_squared_error, r2_score)
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.preprocessing import MinMaxScaler

df = pd.read_csv('heart.csv')
df.fillna(df.mean(), inplace=True)
X = df.drop(columns=['TenYearCHD'])
target = df['TenYearCHD']

X_train, X_test, y_train, y_test = train_test_split(X, target, test_size=0.2, random_state=42)
print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)
print("X_test shape:", X_test.shape)
print("y_test shape:", y_test.shape)

print("\n==== Logistic Regression =====")
LR_model = make_pipeline(
    SimpleImputer(strategy='mean'),
    MinMaxScaler(),
    LogisticRegression()
)
LR_model.fit(X_train, y_train)

y_training_pred = LR_model.predict(X_train)
y_testing_pred = LR_model.predict(X_test)

training_acc = accuracy_score(y_train, y_training_pred)
testing_acc = accuracy_score(y_test, y_testing_pred)
print(f"Logistic Regression - Training accuracy : {training_acc:.4f}")
print(f"Logistic Regression - Testing accuracy : {testing_acc:.4f}")
```

```

ConfusionMatrixDisplay.from_estimator(LR_model, X_test, y_test)
plt.title("Logistic Regression - Confusion Matrix")
plt.show()

print("Logistic Regression - Classification Report:\n", classification_report(y_test, y_testing_pred))

y_test_prob = LR_model.predict_proba(X_test)[: , 1]
fpr, tpr, _ = roc_curve(y_test, y_test_prob)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Logistic Regression - ROC Curve')
plt.legend(loc="lower right")
plt.show()

print("\n==== Multiple Linear Regression =====")
MLR_model = make_pipeline(
    SimpleImputer(strategy='mean'),
    MinMaxScaler(),
    LinearRegression()
)
MLR_model.fit(X_train, y_train)

y_train_pred_mlr = MLR_model.predict(X_train)
y_test_pred_mlr = MLR_model.predict(X_test)

train_mse = mean_squared_error(y_train, y_train_pred_mlr)
test_mse = mean_squared_error(y_test, y_test_pred_mlr)
train_r2 = r2_score(y_train, y_train_pred_mlr)
test_r2 = r2_score(y_test, y_test_pred_mlr)
print(f"Multiple Linear Regression - Training MSE: {train_mse:.4f}")
print(f"Multiple Linear Regression - Testing MSE: {test_mse:.4f}")

```



```

print(f"Multiple Linear Regression - Testing MSE: {test_mse:.4f}")
print(f"Multiple Linear Regression - Training R^2 Score: {train_r2:.4f}")
print(f"Multiple Linear Regression - Testing R^2 Score: {test_r2:.4f}")

y_train_pred_mlr_rounded = np.round(y_train_pred_mlr)
y_test_pred_mlr_rounded = np.round(y_test_pred_mlr)

mlr_training_acc = accuracy_score(y_train, y_train_pred_mlr_rounded)
mlr_testing_acc = accuracy_score(y_test, y_test_pred_mlr_rounded)
print(f"Multiple Linear Regression (Rounded) - Training accuracy: {mlr_training_acc:.4f}")
print(f"Multiple Linear Regression (Rounded) - Testing accuracy: {mlr_testing_acc:.4f}")

ConfusionMatrixDisplay.from_predictions(y_test, y_test_pred_mlr_rounded)
plt.title("Multiple Linear Regression (Rounded) - Confusion Matrix")
plt.show()

print("Multiple Linear Regression (Rounded) - Classification Report:\n",
      classification_report(y_test, y_test_pred_mlr_rounded))

fpr_mlr, tpr_mlr, _ = roc_curve(y_test, y_test_pred_mlr_rounded)
roc_auc_mlr = auc(fpr_mlr, tpr_mlr)

plt.figure()
plt.plot(fpr_mlr, tpr_mlr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc_mlr:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Multiple Linear Regression (Rounded) - ROC Curve')
plt.legend(loc="lower right")
plt.show()

predict_df = pd.read_csv("heart2.csv", skiprows=2)
predict_df = predict_df.iloc[:, 1:]
predict_df = predict_df.iloc[1:].reset_index(drop=True)
predict_df = predict_df[X_train.columns]

predictions_LR = LR_model.predict(predict_df)

```

```

predict_df = pd.read_csv("heart2.csv", skiprows=2)
predict_df = predict_df.iloc[:, 1:]
predict_df = predict_df.iloc[1:].reset_index(drop=True)
predict_df = predict_df[X_train.columns]

predictions_LR = LR_model.predict(predict_df)
predictions_MLR_rounded = np.round(MLR_model.predict(predict_df))

predict_df['TenYearCHD_Logistic'] = predictions_LR
predict_df['TenYearCHD_Linear'] = predictions_MLR_rounded

print("\nPredictions (Logistic Regression):", predictions_LR)
print("Predictions (Multiple Linear Regression Rounded):", predictions_MLR_rounded)


```

```
X_train shape: (3390, 15)
y_train shape: (3390,)
X_test shape: (848, 15)
y_test shape: (848,)
```

```
==== Logistic Regression ====
```

```
Logistic Regression - Training accuracy : 0.8546
```


```
Logistic Regression - Testing accuracy : 0.8573
```




```
Logistic Regression - Classification Report:
              precision    recall  f1-score   support

     0       0.86         0.99         0.92         724
     1       0.60         0.07         0.13         124

 accuracy          0.86         848
 macro avg         0.73         0.53         0.53         848
 weighted avg      0.82         0.86         0.81         848
```




```

===== Multiple Linear Regression =====
Multiple Linear Regression - Training MSE: 0.1167
Multiple Linear Regression - Testing MSE: 0.1160
Multiple Linear Regression - Training R^2 Score: 0.1017
Multiple Linear Regression - Testing R^2 Score: 0.0708
Multiple Linear Regression (Rounded) - Training accuracy: 0.8513
Multiple Linear Regression (Rounded) - Testing accuracy: 0.8550

```

Multiple Linear Regression (Rounded) - Confusion Matrix




Multiple Linear Regression (Rounded) - Classification Report:

	precision	recall	f1-score	support
0	0.86	1.00	0.92	724
1	0.60	0.02	0.05	124
accuracy			0.85	848
macro avg	0.73	0.51	0.48	848
weighted avg	0.82	0.85	0.79	848

Multiple Linear Regression (Rounded) - Classification Report:

	precision	recall	f1-score	support
0	0.86	1.00	0.92	724
1	0.60	0.02	0.05	124
accuracy			0.85	848
macro avg	0.73	0.51	0.48	848
weighted avg	0.82	0.85	0.79	848

Multiple Linear Regression (Rounded) - ROC Curve



Predictions (Logistic Regression): [1 0 1 0 0 1 1 0 1 0]

Predictions (Multiple Linear Regression Rounded): [1. 0. 1. 1. 0. 1. 1. 0. 1. 0.]

Q4. Compute PCA components for the heart disease data. Predict heart disease with the PCA features (consider #PCA features = [1,2,3,4,5]) and evaluate the performance in terms of confusion matrix. Note down your observations

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score


df = pd.read_csv('heart.csv')
df.fillna(df.mean(), inplace=True)
X = df.drop(columns=['TenYearCHD'], axis=1)
y = df['TenYearCHD']
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
accuracies = []
conf_matrices = []
for n_components in range(1, 8):
    pca = PCA(n_components=n_components)
    X_pca = pca.fit_transform(X_scaled)
    X_train, X_test, y_train, y_test = train_test_split(X_pca, y, test_size=0.3, random_state=42)
    model = LogisticRegression()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    conf_matrix = confusion_matrix(y_test, y_pred)
    accuracy = accuracy_score(y_test, y_pred)
    conf_matrices.append(conf_matrix)
    accuracies.append(accuracy)
    print(f"PCA Components: {n_components}")
    print("Confusion Matrix:")
    print(conf_matrix)
    print(f"Accuracy: {accuracy:.5f}\n")

plt.plot(range(1, 8), accuracies, marker='o')
plt.xlabel('Number of PCA Components')
plt.ylabel('Accuracy')
plt.title('Accuracy vs. Number of PCA Components')
plt.grid(True)
plt.show()


predict_df = pd.read_csv("heart2.csv", skiprows=2)
predict_df = predict_df.iloc[:, 1:]
predict_df = predict_df.iloc[1:].reset_index(drop=True)
predict_df_scaled = scaler.transform(predict_df)

for n_components in range(1, 8):
    pca = PCA(n_components=n_components)
    predict_df_pca = pca.fit_transform(predict_df_scaled)
    X_pca = PCA(n_components=n_components).fit_transform(X_scaled)
    X_train, X_test, y_train, y_test = train_test_split(X_pca, y, test_size=0.3, random_state=42)
    model = LogisticRegression()
    model.fit(X_train, y_train)
    predictions = model.predict(predict_df_pca)
    print(f"Predictions with {n_components} PCA Components:")
    print(predictions)
```


PCA Components: 1
Accuracy: 0.85456




PCA Components: 2
Accuracy: 0.85377




PCA Components: 3
Accuracy: 0.85299




PCA Components: 4
Accuracy: 0.85456





PCA Components: 5
Accuracy: 0.85377



PCA Components: 6
Accuracy: 0.85535



PCA Components: 7
Accuracy: 0.85299




```
Predictions with 1 PCA Components:
[0 0 1 0 0 1 1 0 0 0]
Predictions with 2 PCA Components:
[0 0 1 0 0 1 1 0 0 0]
Predictions with 3 PCA Components:
[0 0 1 0 0 1 1 0 0 0]
Predictions with 4 PCA Components:
[0 0 1 0 0 1 1 0 0 0]
Predictions with 5 PCA Components:
[0 0 1 0 0 1 1 0 0 0]
Predictions with 6 PCA Components:
[0 0 1 0 0 1 1 0 0 0]
Predictions with 7 PCA Components:
[0 0 1 0 0 1 1 0 0 0]
```

Observations

PCA is used to reduce the dimensionality of the dataset while retaining as much variance as possible. This can help in improving the performance and reducing the complexity of the model.

Each principal component will explain a certain amount of the variance in the dataset. The first few components usually capture the majority of the variance.

As we increase the number of PCA components, we observe changes in the accuracy of the logistic regression model. Initially, with fewer components, the model doesn't perform well due to loss of information. As we add more components, the performance will improve up to a certain point and then plateau or even decrease due to overfitting.

The confusion matrices shows how well the model is performing in terms of true positives, true negatives, false positives, and false negatives. We observe that with more PCA components, the confusion matrix becomes more balanced, indicating better model performance.

Plotting the accuracy against the number of PCA components will help we visualize the optimal number of components. We find that after a certain number of components, the accuracy does not improve significantly.

The predictions on the new dataset using different numbers of PCA components shows how well the model generalizes to unseen data. We observe differences in predictions with different numbers of components.

Q5. Load Cameraman image from python libraries and reduce the dimensionality using SVD, Check its visual appearances (original versus new image) for the different numbers of SVD components. Additionally, find the correlation between the original and reconstructed images from the different numbers of SVD components (say 1,2,...,9,10)


```
import numpy as np
from skimage import data
import matplotlib.pyplot as plt
from numpy.linalg import svd

image = data.camera()
U, s, Vt = svd(image, full_matrices=False)
def reconstruct_image(U, s, Vt, k):
    return np.matrix(U[:, :k]) * np.diag(s[:k]) * np.matrix(Vt[:, :k, :])
plt.figure(figsize=(15, 10))
plt.subplot(3, 4, 1)
plt.imshow(image, cmap='gray')
plt.title('Original Image')
plt.axis('off')
correlations = []
for i, k in enumerate([1, 2, 3, 4, 5, 6, 7, 8, 9, 10]):
    reconstructed = reconstruct_image(U, s, Vt, k)
    correlation = np.corrcoef(image.flatten(), np.array(reconstructed).flatten())[0,1]
    correlations.append(correlation)
    plt.subplot(3, 4, i+2)
    plt.imshow(reconstructed, cmap='gray')
    plt.title(f'k={k}\nCorr={correlation:.3f}')
    plt.axis('off')
plt.figure(figsize=(10, 5))
plt.plot(range(1, 11), correlations, 'ro-')
plt.xlabel('Number of SVD Components')
plt.ylabel('Correlation Coefficient')
plt.title('Correlation between Original and Reconstructed Images')
plt.grid(True)
plt.show()
print("\nCorrelation coefficients:")
for k, corr in enumerate(correlations, 1):
    print(f"Components: {k}, Correlation: {corr:.4f}")
```


Original Image



k=1
Corr=0.689



k=2
Corr=0.822



k=3
Corr=0.895



k=4
Corr=0.925



k=5
Corr=0.938



k=6
Corr=0.945



k=7
Corr=0.950



k=8
Corr=0.955

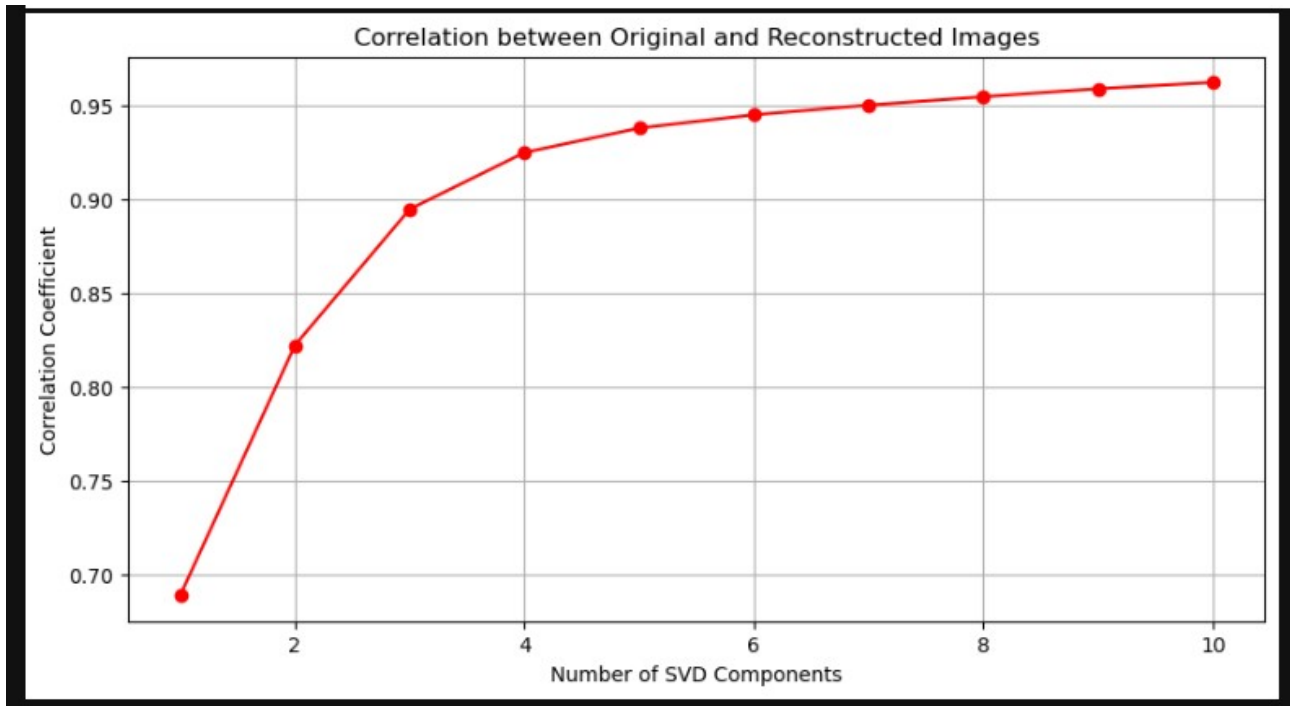


k=9
Corr=0.959



k=10
Corr=0.962





```
Correlation coefficients:  
Components: 1, Correlation: 0.6888  
Components: 2, Correlation: 0.8221  
Components: 3, Correlation: 0.8946  
Components: 4, Correlation: 0.9248  
Components: 5, Correlation: 0.9378  
Components: 6, Correlation: 0.9449  
Components: 7, Correlation: 0.9501  
Components: 8, Correlation: 0.9545  
Components: 9, Correlation: 0.9588  
Components: 10, Correlation: 0.9622
```

Q6. Load Cocktail Party Problem dataset from kaggle to perform ICA on separating the audios of different speakers. Test PCA and compare its performance with ICA's in source separation problem.

<https://www.kaggle.com/datasets/anashamoutni/cocktail-party-problem-cities-of-the-us>

```

import numpy as np
import matplotlib.pyplot as plt
import librosa
from sklearn.decomposition import FastICA, PCA
from sklearn.preprocessing import StandardScaler
from scipy.signal import butter, filtfilt
from IPython.display import Audio
from IPython.display import Audio, display

original_paths = [
    "Washington - American Woman.mp3",
    "New York - American Man.mp3",
    "Miami - Australian Woman.mp3",
    "Chicago - British Woman.mp3"
]

original_audios = [librosa.load(path, sr=44100)[0][:44100 * 60] for path in original_paths]

for i, audio in enumerate(original_audios):
    plt.subplot(4, 1, i + 1)
    plt.plot(audio)
    plt.title(f'Original Audio {i+1}')
plt.tight_layout()
plt.show()

for audio in original_audios:
    display(Audio(audio, rate=44100))

mixed_path = "Audio mix.mp3"
mixed_audio, _ = librosa.load(mixed_path, sr=44100)
mixed_audio = mixed_audio[:44100 * 60]

plt.plot(mixed_audio)
plt.title('Mixed Audio')
plt.show()

display(Audio(mixed_audio, rate=44100))

def play_audio_in_jupyter(audio_results):
    """
    Play audio results specifically in Jupyter Notebook
    """
    print("Original Sources:")
    for i, source in enumerate(audio_results['original'], 1):
        print(f'Original Source {i}')
        display(source)

    print("\nMixed Audio:")
    display(audio_results['mixed'])

    print("\nICA Separated Sources:")
    for i, source in enumerate(audio_results['ica_separated'], 1):
        print(f'ICA Separated Source {i}')
        display(source)

    print("\nPCA Separated Sources:")
    for i, source in enumerate(audio_results['pca_separated'], 1):
        print(f'PCA Separated Source {i}')
        display(source)

```

```

def load_and_preprocess_audio(file_paths, duration=60, sr=44100):
    audios = []
    for path in file_paths:
        audio, _ = librosa.load(path, sr=sr)
        # Trim to specified duration
        audio = audio[:sr * duration]
        # Pad if audio is shorter than desired length
        if len(audio) < sr * duration:
            audio = np.pad(audio, (0, sr * duration - len(audio)))
        audios.append(audio)
    return np.array(audios)

def apply_bandpass_filter(audio, lowcut=500, highcut=3000, sr=44100, order=5):
    nyquist = 0.5 * sr
    low = lowcut / nyquist
    high = highcut / nyquist
    b, a = butter(order, [low, high], btype='band')
    return filtfilt(b, a, audio)

def separate_sources(mixed_signals, n_components, method='ica'):
    X = mixed_signals.T
    scaler = StandardScaler()
    X_scaled = scaler.fit_transform(X)

    if method.lower() == 'ica':
        separator = FastICA(n_components=n_components, random_state=42)
    else: # PCA
        separator = PCA(n_components=n_components)

    # Perform separation
    separated = separator.fit_transform(X_scaled)

    return separated.T

def plot_waveforms(signals, titles, figsize=(15, 10)):
    n_signals = len(signals)
    plt.figure(figsize=figsize)

    for i, (signal, title) in enumerate(zip(signals, titles)):
        plt.subplot(n_signals, 1, i + 1)
        plt.plot(signal)
        plt.title(title)
        plt.xlabel('Sample')
        plt.ylabel('Amplitude')

    plt.tight_layout()
    plt.show()

```

```

# Main processing pipeline
def process_audio_separation(original_paths, mixed_path, duration=60, sr=44100):
    original_sources = load_and_preprocess_audio(original_paths, duration, sr)

    mixed_audio, _ = librosa.load(mixed_path, sr=sr)
    mixed_audio = mixed_audio[:sr * duration]

    n_sources = len(original_paths)
    mixing_matrix = np.random.rand(n_sources, n_sources)
    artificial_mixtures = np.dot(mixing_matrix, original_sources)

    all_mixtures = np.vstack([mixed_audio, artificial_mixtures])

    filtered_mixtures = np.array([apply_bandpass_filter(mix) for mix in all_mixtures])

    separated_ica = separate_sources(filtered_mixtures, n_components=n_sources, method='ica')
    separated_pca = separate_sources(filtered_mixtures, n_components=n_sources, method='pca')

    plot_waveforms(original_sources,
                    [f'Original Source {i+1}' for i in range(n_sources)])

    plot_waveforms([mixed_audio], ['Mixed Audio'])



    plot_waveforms(separated_ica,
                    [f'ICA Separated Source {i+1}' for i in range(n_sources)])


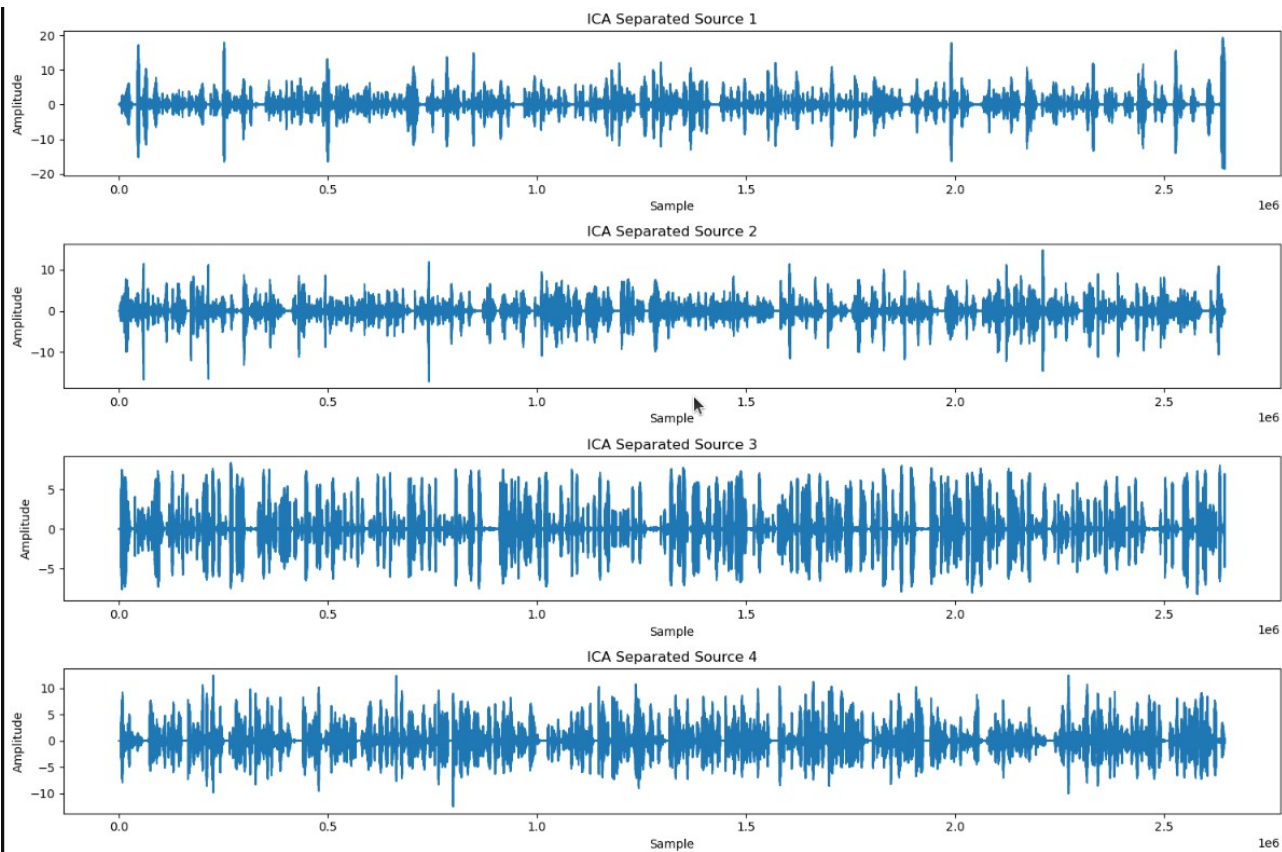
    plot_waveforms(separated_pca,
                    [f'PCA Separated Source {i+1}' for i in range(n_sources)])

    return {
        'original': [Audio(source, rate=sr) for source in original_sources],
        'mixed': Audio(mixed_audio, rate=sr),
        'ica_separated': [Audio(source, rate=sr) for source in separated_ica],
        'pca_separated': [Audio(source, rate=sr) for source in separated_pca]
    }

results = process_audio_separation(original_paths, mixed_path)
play_audio_in_jupyter(results)

```



Original Sources:
Original Source 1




Original Source 2



Original Source 3



Original Source 4



Mixed Audio:



ICA Separated Sources:
ICA Separated Source 1




ICA Separated Source 2




ICA Separated Source 3



ICA Separated Source 4




PCA Separated Sources:
PCA Separated Source 1



PCA Separated Source 2



PCA Separated Source 3



PCA Separated Source 4



ICA is more effective for source separation in audio signals because it tries to find components that are statistically independent, which aligns well with the nature of mixed audio sources. PCA, on the other hand, focuses on capturing the maximum variance, which might not necessarily correspond to independent sources.